

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
```

```
In [2]: data=pd.read_csv("water_potability.csv")
```

```
In [3]: data.head()
```

Out[3]:

	ph	Hardness	Solids	Chloramines	Sulfate	Conductivity	Organic_carbon
0	NaN	204.890455	20791.318981	7.300212	368.516441	564.308654	10.379783
1	3.716080	129.422921	18630.057858	6.635246	NaN	592.885359	15.180013
2	8.099124	224.236259	19909.541732	9.275884	NaN	418.606213	16.868637
3	8.316766	214.373394	22018.417441	8.059332	356.886136	363.266516	18.436524
4	9.092223	181.101509	17978.986339	6.546600	310.135738	398.410813	11.558279



```
In [4]: data.shape
```

Out[4]: (3276, 10)

```
In [5]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3276 entries, 0 to 3275
Data columns (total 10 columns):
#   Column          Non-Null Count  Dtype  
---  -
0   ph              2785 non-null   float64
1   Hardness        3276 non-null   float64
2   Solids          3276 non-null   float64
3   Chloramines     3276 non-null   float64
4   Sulfate         2495 non-null   float64
5   Conductivity    3276 non-null   float64
6   Organic_carbon  3276 non-null   float64
7   Trihalomethanes 3114 non-null   float64
8   Turbidity       3276 non-null   float64
9   Potability      3276 non-null   int64  
dtypes: float64(9), int64(1)
memory usage: 256.1 KB
```

```
In [6]: data.isnull().sum() #yes we have empty block in the dataset
```

```
Out[6]: ph                491  
Hardness                0  
Solids                  0  
Chloramines             0  
Sulfate                 781  
Conductivity            0  
Organic_carbon          0  
Trihalomethanes        162  
Turbidity               0  
Potability              0  
dtype: int64
```

## Preprocessing

```
In [7]: # total rows=3276  
data.count()  
# ph,Sulfate,Trihalomethanes have less value
```

```
Out[7]: ph                2785  
Hardness                3276  
Solids                  3276  
Chloramines             3276  
Sulfate                 2495  
Conductivity            3276  
Organic_carbon          3276  
Trihalomethanes        3114  
Turbidity               3276  
Potability              3276  
dtype: int64
```

```
In [8]: data.mean()
```

```
Out[8]: ph                7.080795  
Hardness                196.369496  
Solids                  22014.092526  
Chloramines             7.122277  
Sulfate                 333.775777  
Conductivity            426.205111  
Organic_carbon          14.284970  
Trihalomethanes        66.396293  
Turbidity               3.966786  
Potability              0.390110  
dtype: float64
```

```
In [9]: data.describe()
```

```
Out[9]:
```

	ph	Hardness	Solids	Chloramines	Sulfate	Conductivity	Organic
count	2785.000000	3276.000000	3276.000000	3276.000000	2495.000000	3276.000000	3276.000000
mean	7.080795	196.369496	22014.092526	7.122277	333.775777	426.205111	16.180011
std	1.594320	32.879761	8768.570828	1.583085	41.416840	80.824064	3.379786
min	0.000000	47.432000	320.942611	0.352000	129.000000	181.483754	10.379786
25%	6.093092	176.850538	15666.690297	6.127421	307.699498	365.734414	13.894411
50%	7.036752	196.967627	20927.833607	7.130299	333.073546	421.884968	15.180011
75%	8.062066	216.667456	27332.762127	8.114887	359.950170	481.792304	16.868633
max	14.000000	323.124000	61227.196008	13.127000	481.030642	753.342620	20.140322



```
In [10]: data.fillna(data.mean(),inplace=True)
```

```
In [11]: data #you can find empty blocks are filled with mean of respective columns
```

```
Out[11]:
```

	ph	Hardness	Solids	Chloramines	Sulfate	Conductivity	Organic_carbo
0	7.080795	204.890455	20791.318981	7.300212	368.516441	564.308654	10.379786
1	3.716080	129.422921	18630.057858	6.635246	333.775777	592.885359	15.180011
2	8.099124	224.236259	19909.541732	9.275884	333.775777	418.606213	16.868633
3	8.316766	214.373394	22018.417441	8.059332	356.886136	363.266516	18.436521
4	9.092223	181.101509	17978.986339	6.546600	310.135738	398.410813	11.558271
...	...	...	...	...	...	...	...
3271	4.668102	193.681735	47580.991603	7.166639	359.948574	526.424171	13.894411
3272	7.808856	193.553212	17329.802160	8.061362	333.775777	392.449580	19.903221
3273	9.419510	175.762646	33155.578218	7.350233	333.775777	432.044783	11.039071
3274	5.126763	230.603758	11983.869376	6.303357	333.775777	402.883113	11.168941
3275	7.874671	195.102299	17404.177061	7.509306	333.775777	327.459760	16.140322

3276 rows × 10 columns



```
In [12]: # now check the null value
data.isnull().sum()
#you can see none of column has empty value
```

```
Out[12]: ph                0
Hardness                  0
Solids                    0
Chloramines               0
Sulfate                   0
Conductivity              0
Organic_carbon            0
Trihalomethanes           0
Turbidity                 0
Potability                0
dtype: int64
```


```
In [13]: # checking for the duplicate value
data.duplicated().sum() #it means no duplicate values are present
```

```
Out[13]: 0
```

```
In [14]: # Observing the nature of data distribution
for i in data.columns:
    sns.distplot(data[i])
    plt.title(i,size=5)
    plt.show()
# we can observe that all data is normally distributed, So we can use it direct
```


C:\Users\A to Z Infosys\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)




C:\Users\A to Z Infosys\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)




```
C:\Users\A to Z Infosys\anaconda3\lib\site-packages\seaborn\distributions.py:
2619: FutureWarning: `distplot` is a deprecated function and will be removed
in a future version. Please adapt your code to use either `displot` (a figure
-level function with similar flexibility) or `histplot` (an axes-level functi
on for histograms).
```

```
warnings.warn(msg, FutureWarning)
```




```
C:\Users\A to Z Infosys\anaconda3\lib\site-packages\seaborn\distributions.py:
2619: FutureWarning: `distplot` is a deprecated function and will be removed
in a future version. Please adapt your code to use either `displot` (a figure
-level function with similar flexibility) or `histplot` (an axes-level functi
on for histograms).
      warnings.warn(msg, FutureWarning)
```



```
C:\Users\A to Z Infosys\anaconda3\lib\site-packages\seaborn\distributions.py:
2619: FutureWarning: `distplot` is a deprecated function and will be removed
in a future version. Please adapt your code to use either `displot` (a figure
-level function with similar flexibility) or `histplot` (an axes-level functi
on for histograms).
```


```
warnings.warn(msg, FutureWarning)
```





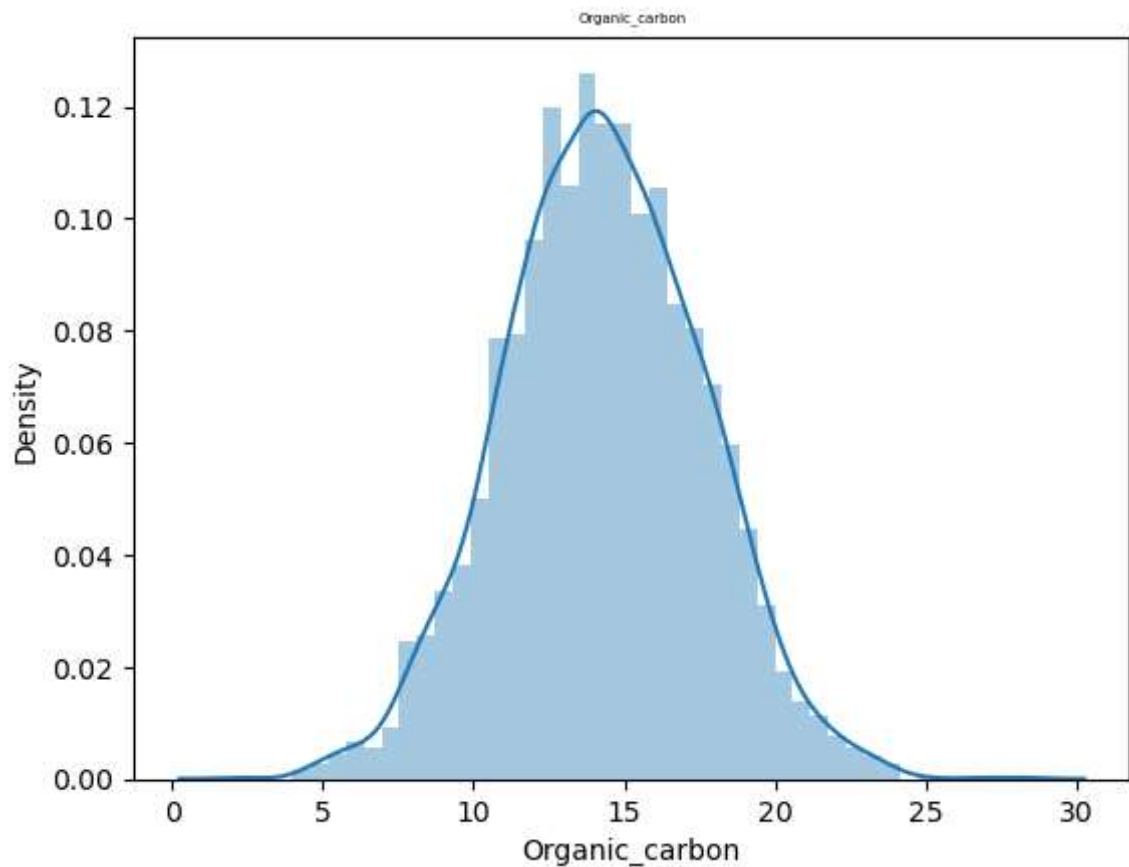
```
C:\Users\A to Z Infosys\anaconda3\lib\site-packages\seaborn\distributions.py:
2619: FutureWarning: `distplot` is a deprecated function and will be removed
in a future version. Please adapt your code to use either `displot` (a figure
-level function with similar flexibility) or `histplot` (an axes-level functi
on for histograms).
```

```
warnings.warn(msg, FutureWarning)
```



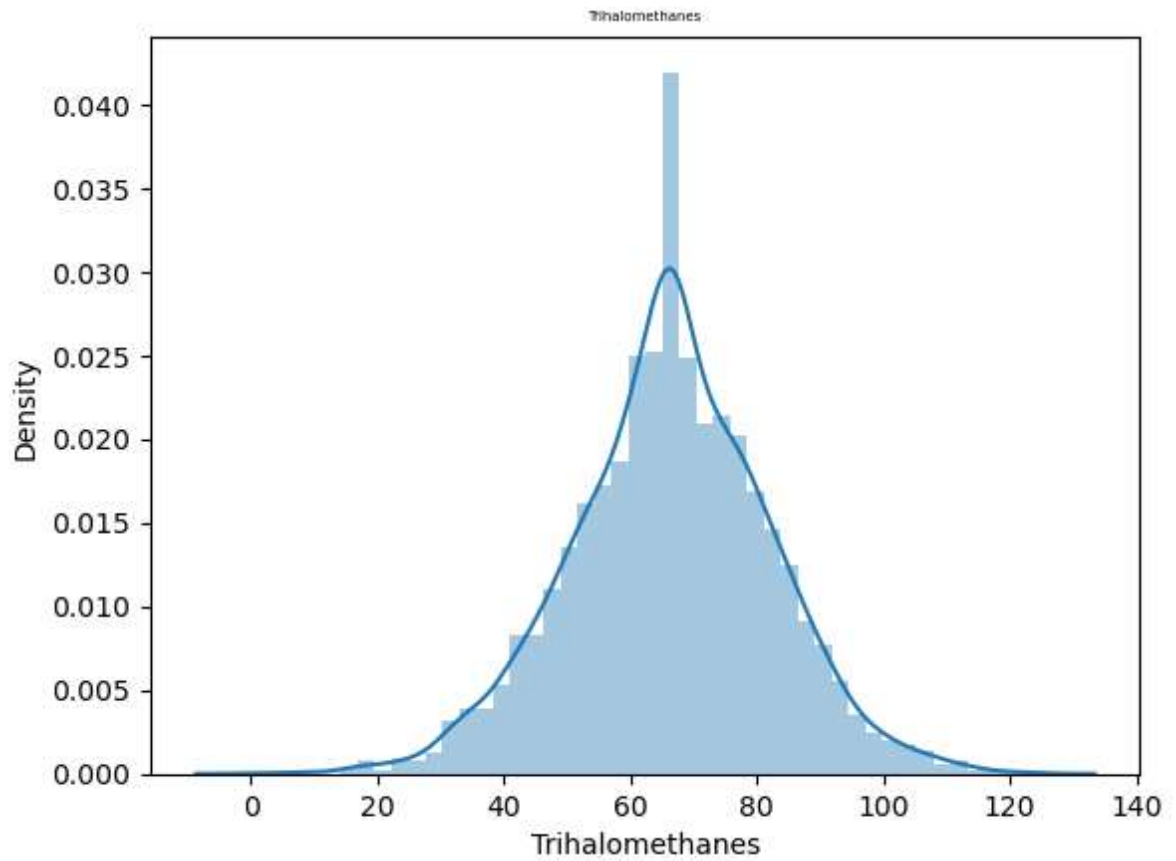
```
C:\Users\A to Z Infosys\anaconda3\lib\site-packages\seaborn\distributions.py:
2619: FutureWarning: `distplot` is a deprecated function and will be removed
in a future version. Please adapt your code to use either `displot` (a figure
-level function with similar flexibility) or `histplot` (an axes-level functi
on for histograms).
```

```
warnings.warn(msg, FutureWarning)
```




```
C:\Users\A to Z Infosys\anaconda3\lib\site-packages\seaborn\distributions.py:
2619: FutureWarning: `distplot` is a deprecated function and will be removed
in a future version. Please adapt your code to use either `displot` (a figure
-level function with similar flexibility) or `histplot` (an axes-level functi
on for histograms).
```

```
warnings.warn(msg, FutureWarning)
```




```
C:\Users\A to Z Infosys\anaconda3\lib\site-packages\seaborn\distributions.py:
2619: FutureWarning: `distplot` is a deprecated function and will be removed
in a future version. Please adapt your code to use either `displot` (a figure
-level function with similar flexibility) or `histplot` (an axes-level functi
on for histograms).
```

```
warnings.warn(msg, FutureWarning)
```




```
C:\Users\A to Z Infosys\anaconda3\lib\site-packages\seaborn\distributions.py:
2619: FutureWarning: `distplot` is a deprecated function and will be removed
in a future version. Please adapt your code to use either `displot` (a figure
-level function with similar flexibility) or `histplot` (an axes-level functi
on for histograms).
```

```
warnings.warn(msg, FutureWarning)
```



```
In [15]: # checking the correlation of all the parameter,  
# if correlation value is high than we can keep any one attribute  
sns.heatmap(data.corr())  
# most of the colors are darker means our data is less corelated with each oth
```


Out[15]: <AxesSubplot:>



```
In [16]: # Now checking the distribution of target column  
data['Potability'].value_counts()
```

Out[16]: 0 1998  
1 1278  
Name: Potability, dtype: int64

```
In [17]: data['Potability'].value_counts().plot(kind='bar')
plt.xticks(rotation='vertical')
plt.show()
```



```
In [18]: x=data.drop(columns=['Potability'])
```

```
In [19]: y=data['Potability']
```

```
In [20]: # scaler=StandardScaler()
# x=scaler.fit_transform(x)
```



In [21]: x

Out[21]:

	ph	Hardness	Solids	Chloramines	Sulfate	Conductivity	Organic_carbo
0	7.080795	204.890455	20791.318981	7.300212	368.516441	564.308654	10.37976
1	3.716080	129.422921	18630.057858	6.635246	333.775777	592.885359	15.18001
2	8.099124	224.236259	19909.541732	9.275884	333.775777	418.606213	16.86863
3	8.316766	214.373394	22018.417441	8.059332	356.886136	363.266516	18.43652
4	9.092223	181.101509	17978.986339	6.546600	310.135738	398.410813	11.55827
...	...	...	...	...	...	...	...
3271	4.668102	193.681735	47580.991603	7.166639	359.948574	526.424171	13.89441
3272	7.808856	193.553212	17329.802160	8.061362	333.775777	392.449580	19.90322
3273	9.419510	175.762646	33155.578218	7.350233	333.775777	432.044783	11.03907
3274	5.126763	230.603758	11983.869376	6.303357	333.775777	402.883113	11.16894
3275	7.874671	195.102299	17404.177061	7.509306	333.775777	327.459760	16.14036

3276 rows × 9 columns



## Resampling of data (optinal)

In [22]: y.value\_counts()

Out[22]: 0 1998  
1 1278  
Name: Potability, dtype: int64

```
In [23]: # Now we are doing data oversampling- means it make the equal number of sample  
# this is optional step  
from imblearn import over_sampling  
from imblearn.over_sampling import RandomOverSampler
```

```
In [26]: ros=RandomOverSampler(random_state=0)  
x_resampled,y_resampled=ros.fit_resample(x,y)
```

```
In [27]: # x_resampled['Potability'].value_counts()
```

```
In [28]: x_resampled
```

```
Out[28]:
```

	ph	Hardness	Solids	Chloramines	Sulfate	Conductivity	Organic_carbo
0	7.080795	204.890455	20791.318981	7.300212	368.516441	564.308654	10.37978
1	3.716080	129.422921	18630.057858	6.635246	333.775777	592.885359	15.18001
2	8.099124	224.236259	19909.541732	9.275884	333.775777	418.606213	16.86863
3	8.316766	214.373394	22018.417441	8.059332	356.886136	363.266516	18.43652
4	9.092223	181.101509	17978.986339	6.546600	310.135738	398.410813	11.55827
...	...	...	...	...	...	...	...
3991	5.913755	175.326062	12044.624691	8.368785	347.880372	380.967166	12.53082
3992	7.672910	152.878305	22989.351184	6.231913	343.439017	401.140879	14.25266
3993	7.080795	129.883297	8906.975623	6.827901	327.551789	525.224717	17.65747
3994	7.226593	207.832229	22097.413223	5.862717	354.708382	353.654046	10.40105
3995	7.080795	182.317256	30430.211752	6.151217	350.448584	479.957076	16.57656

3996 rows × 9 columns



```
In [29]: y_resampled
```

```
Out[29]:
```

```
0      0
1      0
2      0
3      0
4      0
..
3991    1
3992    1
3993    1
3994    1
3995    1
Name: Potability, Length: 3996, dtype: int64
```

```
In [30]: r_data=pd.concat([x_resampled,y_resampled],axis=1) #our data is ready with res
```

```
In [31]: # r_data.duplicated().sum()
# as it adds the duplicate value we won't use sampling
```

## Data splitting and Model Trainig

```
In [32]: # Now it time to start Model training and we are using four models
# i)Random Forest
# ii)DecisionTree
# iii)AdaBoost
# iv)Naive Bayes
#v) Logistic Regression
#vi) K Nearest Neighbour
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, confusion_matrix

# models
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.neighbors import KNeighborsClassifier
```

```
In [33]: All_model=[

    ("Random Forest", RandomForestClassifier(random_state=10)),
#    ("Linear Regression", LinearRegression()),
    ("Decision Tree", DecisionTreeClassifier(random_state=42)),
    ("Naive Bayes", GaussianNB()),
    ("Ada Boost", AdaBoostClassifier(random_state=42)),
    ("K Neighbour Classifier", KNeighborsClassifier(n_neighbors=2))
]
```

```
In [45]: # X=data.drop(columns='Potability')
X=r_data.drop(columns='Potability')
# Y=data['Potability']
Y=r_data['Potability']
```

In [46]: X

Out[46]:

	ph	Hardness	Solids	Chloramines	Sulfate	Conductivity	Organic_carbo
0	7.080795	204.890455	20791.318981	7.300212	368.516441	564.308654	10.37978
1	3.716080	129.422921	18630.057858	6.635246	333.775777	592.885359	15.18001
2	8.099124	224.236259	19909.541732	9.275884	333.775777	418.606213	16.86863
3	8.316766	214.373394	22018.417441	8.059332	356.886136	363.266516	18.43652
4	9.092223	181.101509	17978.986339	6.546600	310.135738	398.410813	11.55827
...	...	...	...	...	...	...	...
3991	5.913755	175.326062	12044.624691	8.368785	347.880372	380.967166	12.53082
3992	7.672910	152.878305	22989.351184	6.231913	343.439017	401.140879	14.25266
3993	7.080795	129.883297	8906.975623	6.827901	327.551789	525.224717	17.65747
3994	7.226593	207.832229	22097.413223	5.862717	354.708382	353.654046	10.40105
3995	7.080795	182.317256	30430.211752	6.151217	350.448584	479.957076	16.57656

3996 rows × 9 columns



In [47]: X\_train, X\_test, Y\_train, Y\_test = train\_test\_split(X, Y, test\_size=0.2, random\_state=42)

In [48]: *# Creating new dataframe to store the accuracy wrt to model used*  
accuracy\_table=pd.DataFrame(columns=["Algo", "Accuracy"])

```
In [49]: for i,j in All_model:
          print(j)
          print(i)
          j.fit(X_train,Y_train)
          predictions=j.predict(X_test)
          acc=accuracy_score(Y_test, predictions)
          temp= {"Algo": i, "Accuracy": acc}
          accuracy_table=accuracy_table.append(temp, ignore_index=True)
```

```
RandomForestClassifier(random_state=10)
Random Forest
```

```
C:\Users\A to Z Infosys\AppData\Local\Temp\ipykernel_14280\1411955508.py:8: FutureWarning: The frame.append method is deprecated and will be removed from pandas in a future version. Use pandas.concat instead.
```

```
accuracy_table=accuracy_table.append(temp, ignore_index=True)
```

```
C:\Users\A to Z Infosys\AppData\Local\Temp\ipykernel_14280\1411955508.py:8: FutureWarning: The frame.append method is deprecated and will be removed from pandas in a future version. Use pandas.concat instead.
```

```
accuracy_table=accuracy_table.append(temp, ignore_index=True)
```

```
C:\Users\A to Z Infosys\AppData\Local\Temp\ipykernel_14280\1411955508.py:8: FutureWarning: The frame.append method is deprecated and will be removed from pandas in a future version. Use pandas.concat instead.
```

```
accuracy_table=accuracy_table.append(temp, ignore_index=True)
```

```
DecisionTreeClassifier(random_state=42)
```

```
Decision Tree
```

```
GaussianNB()
```

```
Naive Bayed
```

```
AdaBoostClassifier(random_state=42)
```

```
Ada Boost
```

```
KNeighborsClassifier(n_neighbors=2)
```

```
K Neighbour Classifier
```

```
C:\Users\A to Z Infosys\AppData\Local\Temp\ipykernel_14280\1411955508.py:8: FutureWarning: The frame.append method is deprecated and will be removed from pandas in a future version. Use pandas.concat instead.
```

```
accuracy_table=accuracy_table.append(temp, ignore_index=True)
```

```
C:\Users\A to Z Infosys\AppData\Local\Temp\ipykernel_14280\1411955508.py:8: FutureWarning: The frame.append method is deprecated and will be removed from pandas in a future version. Use pandas.concat instead.
```

```
accuracy_table=accuracy_table.append(temp, ignore_index=True)
```

```
In [50]: accuracy_table.sort_values(by='Accuracy')
```

Out[50]:

	Algo	Accuracy
2	Naive Bayed	0.55875
3	Ada Boost	0.55875
4	K Neighbour Classifier	0.60125
1	Decision Tree	0.70125
0	Random Forest	0.7775

```
In [51]: # if do the same using StanderScaler
```

```
scaler = StandardScaler()  
X = scaler.fit_transform(X)
```

```
In [52]: X
```

```
Out[52]: array([[ -0.00391605,  0.25956782, -0.15738462, ..., -1.18098616,  
                  1.28842112, -1.28102328],  
                [-2.33332377, -2.01696775, -0.40100535, ...,  0.27773111,  
                 -0.63944346,  0.67326663],  
                [ 0.70107841,  0.8431487 , -0.25677993, ...,  0.79087833,  
                 -0.00497148, -1.16307046],  
                ...,  
                [-0.00391605, -2.00308016, -1.49700636, ...,  1.03057479,  
                 -2.1574769 , -0.16378381],  
                [ 0.09702074,  0.34830866, -0.01015964, ..., -1.17452155,  
                 0.022948 ,  0.46539736],  
                [-0.00391605, -0.4213699 ,  0.92912642, ...,  0.70212205,  
                 -0.49193949,  0.71834676]])
```

```
In [42]: X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_state=42)

accuracy_table2=pd.DataFrame(columns=["Algo2", "Accuracy2"])

for i,j in All_model:
    print(j)
    print(i)
    j.fit(X_train,Y_train)
    predictions=j.predict(X_test)
    acc=accuracy_score(Y_test, predictions)
    temp= {"Algo2": i, "Accuracy2": acc}
    accuracy_table2=accuracy_table2.append(temp, ignore_index=True)
```

RandomForestClassifier(random\_state=10)

Random Forest

C:\Users\A to Z Infosys\AppData\Local\Temp\ipykernel\_14280\2942780521.py:12:  
FutureWarning: The frame.append method is deprecated and will be removed from pandas in a future version. Use pandas.concat instead.

accuracy\_table2=accuracy\_table2.append(temp, ignore\_index=True)

C:\Users\A to Z Infosys\AppData\Local\Temp\ipykernel\_14280\2942780521.py:12:  
FutureWarning: The frame.append method is deprecated and will be removed from pandas in a future version. Use pandas.concat instead.

accuracy\_table2=accuracy\_table2.append(temp, ignore\_index=True)

C:\Users\A to Z Infosys\AppData\Local\Temp\ipykernel\_14280\2942780521.py:12:  
FutureWarning: The frame.append method is deprecated and will be removed from pandas in a future version. Use pandas.concat instead.

accuracy\_table2=accuracy\_table2.append(temp, ignore\_index=True)

DecisionTreeClassifier(random\_state=42)

Decision Tree

GaussianNB()

Naive Bayes

AdaBoostClassifier(random\_state=42)

Ada Boost

KNeighborsClassifier(n\_neighbors=2)

K Neighbour Classifier

C:\Users\A to Z Infosys\AppData\Local\Temp\ipykernel\_14280\2942780521.py:12:  
FutureWarning: The frame.append method is deprecated and will be removed from pandas in a future version. Use pandas.concat instead.

accuracy\_table2=accuracy\_table2.append(temp, ignore\_index=True)

C:\Users\A to Z Infosys\AppData\Local\Temp\ipykernel\_14280\2942780521.py:12:  
FutureWarning: The frame.append method is deprecated and will be removed from pandas in a future version. Use pandas.concat instead.

accuracy\_table2=accuracy\_table2.append(temp, ignore\_index=True)

```
In [43]: accuracy_table2.sort_values(by='Accuracy2')
```

Out[43]:

	Algo2	Accuracy2
1	Decision Tree	0.577744
3	Ada Boost	0.620427
2	Naive Bayed	0.631098
4	K Neighbour Classifier	0.660061
0	Random Forest	0.679878

```
In [57]: # Next part we do Hyper Parameter Tuning to improve our efficiency
param_grid_rfc = {"min_samples_split": [2, 3, 10],
                  "min_samples_leaf": [1, 3, 10],
                  "n_estimators": [100, 200, 500],
                  "random_state": [42]}

grid_rfc = GridSearchCV(RandomForestClassifier(), param_grid_rfc, scoring="acc
grid_rfc.fit(X_train, Y_train)
```

-----  
\_RemoteTraceback

Traceback (most recent call las

t)

\_RemoteTraceback:

"""

Traceback (most recent call last):

File "C:\Users\A to Z Infosys\anaconda3\lib\site-packages\joblib\externa  
ls\loky\process\_executor.py", line 428, in \_process\_worker

    r = call\_item()

File "C:\Users\A to Z Infosys\anaconda3\lib\site-packages\joblib\externa  
ls\loky\process\_executor.py", line 275, in \_\_call\_\_

    return self.fn(\*self.args, \*\*self.kwargs)

File "C:\Users\A to Z Infosys\anaconda3\lib\site-packages\joblib\\_\_parall  
el\_backends.py", line 620, in \_\_call\_\_

    return self.func(\*args, \*\*kwargs)

File "C:\Users\A to Z Infosys\anaconda3\lib\site-packages\joblib\paralle  
l.py", line 288, in \_\_call\_\_

    return [func(\*args, \*\*kwargs)

File "C:\Users\A to Z Infosys\anaconda3\lib\site-packages\joblib\paralle

```
In [55]: rfc_params = grid_rfc.best_params_
rfc = RandomForestClassifier(**rfc_params)
rfc.fit(X_train, Y_train)
predictions = rfc.predict(X_test)
score = accuracy_score(Y_test, predictions)
print("Accuracy Score:", score)
```

Accuracy Score: 0.78



In [ ]: